

```

# SMART ENSEMBLE LANGUAGE LOOKUP SYSTEM
# Complete Project Documentation and Code - Updated with Match Rationale Feature

## PROJECT OVERVIEW
This is an advanced language lookup system that uses ensemble retrieval methods to match PDF document content against a

## FEATURES
- Smart Ensemble Retrieval: LLM + BM25 + TF-IDF for candidate selection
- Multiple Similarity Methods: OpenAI, SentenceTransformers, TF-IDF
- LangGraph Multi-Agent Workflow for enhanced analysis
- **NEW: Intelligent Match Rationale** - Explains why each language match was found
- Comprehensive reasoning includes: exact matches, semantic similarity, technical context
- No vector database dependency (more reliable)
- Professional Streamlit interface with rationale display
- Export capabilities (JSON, Markdown) with match reasoning

## PROJECT STRUCTURE
```
langchain-lookup-project/
├── src/
│ ├── ensemble_retriever.py # Multi-method retrieval system with rationale
│ ├── smart_similarity_search.py # Main search pipeline with reasoning
│ ├── pdf_processor.py # LangChain PDF processing
│ └── langgraph_workflow.py # Multi-agent AI workflow
├── data/
│ ├── languages.csv # AI/ML language database
│ ├── ai_topics.txt # Sample document content
│ └── vectorstore/ # Auto-created directory
├── app.py # Streamlit interface with rationale display
├── requirements.txt # Dependencies
├── .env # Environment variables
└── README.md # Project documentation
```

## REQUIREMENTS (requirements.txt)
```
streamlit
langchain
langchain-openai
langchain-community
langgraph
rank_bm25
scikit-learn
openai
pandas
numpy
pypdf
python-dotenv
tiktoken
sentence-transformers

pydantic
```

## ENVIRONMENT VARIABLES (.env)
```
OPENAI_API_KEY=your_openai_api_key_here
CHROMA_INDEX_PATH=./vectorstore/language_index
LANGCHAIN_TRACING_V2=false
LANGCHAIN_API_KEY=your_langsmith_api_key_here
```

## SOURCE CODE FILES

### 1. ENSEMBLE RETRIEVER (src/ensemble_retriever.py) - Updated with Rationale
```python
import os
import pandas as pd
from typing import List, Dict, Any, Optional
from langchain_openai import ChatOpenAI, OpenAIEmbeddings
from langchain.prompts import PromptTemplate

```

```

from langchain.schema import HumanMessage
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics.pairwise import cosine_similarity
from sentence_transformers import SentenceTransformer
import numpy as np
from rank_bm25 import BM25Okapi
import json

class EnsembleRetriever:
 def __init__(self, languages_csv_path: str):
 self.languages_df = pd.read_csv(languages_csv_path)
 self.languages = self.languages_df['Language'].tolist()

 # Initialize components (with error handling for API quota)
 try:
 self.llm = ChatOpenAI(
 model="gpt-4o-mini", # Faster and cheaper for retrieval
 temperature=0.1,
 openai_api_key=os.getenv('OPENAI_API_KEY')
)
 self.llm_available = True
 except Exception as e:
 print(f"■■■ OpenAI LLM not available: {e}")
 self.llm = None
 self.llm_available = False

 # Initialize embedding models
 try:
 self.openai_embeddings = OpenAIEmbeddings(
 model="text-embedding-3-small",
 openai_api_key=os.getenv('OPENAI_API_KEY')
)
 self.openai_available = True
 except Exception as e:
 print(f"■■■ OpenAI Embeddings not available: {e}")
 self.openai_embeddings = None
 self.openai_available = False

 self.sentence_model = SentenceTransformer('all-MiniLM-L6-v2')

 # Initialize TF-IDF
 self.tfidf_vectorizer = TfidfVectorizer(
 stop_words='english',
 max_features=5000,
 lowercase=True
)

 # Initialize BM25
 tokenized_languages = [lang.lower().split() for lang in self.languages]
 self.bm25 = BM25Okapi(tokenized_languages)

 # Fit TF-IDF on languages
 self.tfidf_vectorizer.fit(self.languages)

 print(f"■■ Ensemble Retriever initialized with {len(self.languages)} languages")

 def _generate_match_rationale(self, chunk: str, term: str, similarity_score: float,
 sources: List[str], retrieval_method: str) -> str:
 """Generate human-readable rationale for why a language match was found"""

 # Extract key phrases from chunk for analysis
 chunk_lower = chunk.lower()
 term_lower = term.lower()

 # Check for exact matches
 exact_match = term_lower in chunk_lower

 # Check for partial word matches
 term_words = term_lower.split()

```

```

chunk_words = chunk_lower.split()
matching_words = [word for word in term_words if any(word in chunk_word for chunk_word in chunk_words)]

Build rationale based on evidence
rationale_parts = []

if exact_match:
 rationale_parts.append(f"Direct mention of '{term}' found in the text")
elif matching_words:
 rationale_parts.append(f"Related terms found: {' '.join(matching_words)}")

Add similarity context
if similarity_score > 0.8:
 rationale_parts.append("Very high semantic similarity")
elif similarity_score > 0.6:
 rationale_parts.append("High semantic similarity")
elif similarity_score > 0.4:
 rationale_parts.append("Moderate semantic similarity")
else:
 rationale_parts.append("Lower but significant semantic relationship")

Add retrieval method context
source_context = {
 "llm": "LLM identified semantic relevance",
 "llm_fallback": "LLM analysis with fallback matching",
 "bm25": "Strong keyword matching using BM25",
 "tfidf": "Statistical text similarity detected"
}

method_explanations = []
for source in sources:
 if source in source_context:
 method_explanations.append(source_context[source])

if method_explanations:
 rationale_parts.append(f"Detection methods: {' '.join(method_explanations)}")

Add embedding method context
embedding_context = {
 "openai": "using OpenAI embeddings for deep semantic understanding",
 "sentence_transformer": "using SentenceTransformer for semantic analysis",
 "tfidf": "using TF-IDF statistical analysis"
}

if retrieval_method in embedding_context:
 rationale_parts.append(embedding_context[retrieval_method])

Look for technical context indicators
tech_indicators = [
 "algorithm", "model", "training", "learning", "neural", "network",
 "data", "prediction", "classification", "optimization", "artificial",
 "intelligence", "machine", "deep", "analysis", "processing"
]

found_indicators = [indicator for indicator in tech_indicators if indicator in chunk_lower]
if found_indicators:
 rationale_parts.append(f"Technical context detected: {' '.join(found_indicators[:3])}")

Combine rationale parts
if rationale_parts:
 return ". ".join(rationale_parts) + "."
else:
 return f"Match found based on semantic similarity ({similarity_score:.3f}) using {retrieval_method} embedding"

def llm_based_retrieval(self, query: str, top_k: int = 10) -> List[Dict[str, Any]]:
 """Use LLM to identify relevant language terms"""

 if not self.llm_available:
 print("■■ LLM not available, skipping LLM retrieval")
 return []

```

```

Prepare a sample of languages for the prompt (to avoid token limits)
sample_languages = self.languages[:50] # First 50 as examples

retrieval_prompt = PromptTemplate(
 input_variables=["query", "languages"],
 template="""
 Given the following query about AI/ML/Data Science content:
 Query: {query}

 From this list of language terms:
 {languages}

 Identify the TOP {top_k} most relevant language terms that would likely appear in or relate to this query.
 Consider:
 1. Semantic similarity and conceptual relationships
 2. Technical context and domain relevance
 3. Common usage patterns in AI/ML literature

 Return ONLY a JSON list of the most relevant terms, like:
 ["Machine Learning", "Deep Learning", "Neural Networks"]

 Focus on quality over quantity - return only highly relevant matches.
 """.replace("{top_k}", str(top_k))
)

try:
 prompt = retrieval_prompt.format(
 query=query[:500], # Limit query length
 languages=", ".join(sample_languages)
)

 response = self.llm.invoke([HumanMessage(content=prompt)])

 # Parse JSON response
 try:
 relevant_terms = json.loads(response.content)
 if isinstance(relevant_terms, list):
 # Filter to ensure terms exist in our database
 filtered_terms = [term for term in relevant_terms if term in self.languages]

 return [{"term": term, "source": "llm", "score": 1.0} for term in filtered_terms[:top_k]]
 except json.JSONDecodeError:
 # Fallback: extract terms from response text
 response_text = response.content.lower()
 found_terms = []
 for lang in self.languages:
 if lang.lower() in response_text and len(found_terms) < top_k:
 found_terms.append({"term": lang, "source": "llm_fallback", "score": 0.8})
 return found_terms

except Exception as e:
 print(f"LLM retrieval error: {e}")
 return []

return []

def bm25_retrieval(self, query: str, top_k: int = 10) -> List[Dict[str, Any]]:
 """Use BM25 for keyword-based retrieval"""
 tokenized_query = query.lower().split()
 scores = self.bm25.get_scores(tokenized_query)

 # Get top k indices
 top_indices = np.argsort(scores)[::-1][:top_k]

 results = []
 for idx in top_indices:
 if scores[idx] > 0: # Only include non-zero scores
 results.append({
 "term": self.languages[idx],
 "source": "bm25",
 "score": float(scores[idx])
 })

```

```

 })

 return results

def tfidf_retrieval(self, query: str, top_k: int = 10) -> List[Dict[str, Any]]:
 """Use TF-IDF for statistical text retrieval"""
 query_vector = self.tfidf_vectorizer.transform([query])
 language_vectors = self.tfidf_vectorizer.transform(self.languages)

 similarities = cosine_similarity(query_vector, language_vectors)[0]

 # Get top k indices
 top_indices = np.argsort(similarities)[::-1][:top_k]

 results = []
 for idx in top_indices:
 if similarities[idx] > 0.01: # Minimum threshold
 results.append({
 "term": self.languages[idx],

 "source": "tfidf",
 "score": float(similarities[idx])
 })

 return results

def ensemble_retrieval(self, query: str, top_k: int = 15) -> List[Dict[str, Any]]:
 """Combine multiple retrieval methods using ensemble approach"""

 # Get results from different methods
 llm_results = self.llm_based_retrieval(query, top_k=8)
 bm25_results = self.bm25_retrieval(query, top_k=10)
 tfidf_results = self.tfidf_retrieval(query, top_k=10)

 # Combine and deduplicate results
 term_scores = {}

 # Weight different methods
 weights = {
 "llm": 0.5, # Highest weight for semantic understanding
 "llm_fallback": 0.3,
 "bm25": 0.3, # Good for keyword matching
 "tfidf": 0.2 # Statistical baseline
 }

 # Aggregate scores
 for results in [llm_results, bm25_results, tfidf_results]:
 for result in results:
 term = result["term"]
 source = result["source"]
 score = result["score"] * weights.get(source, 0.1)

 if term not in term_scores:
 term_scores[term] = {"score": 0, "sources": []}

 term_scores[term]["score"] += score
 term_scores[term]["sources"].append(source)

 # Sort by combined score and return top k
 sorted_terms = sorted(
 [(term, data["score"], data["sources"]) for term, data in term_scores.items()],
 key=lambda x: x[1],
 reverse=True
)[:top_k]

 ensemble_results = []
 for term, score, sources in sorted_terms:
 ensemble_results.append({
 "term": term,
 "ensemble_score": score,

```

```

 "sources": sources,
 "source_count": len(set(sources))
 })

 return ensemble_results

def compute_final_similarity(self, chunk: str, candidate_terms: List[Dict],
 method: str = "openai") -> List[Dict[str, Any]]:
 """Compute final similarity scores using specified embedding method and generate rationale"""

 if not candidate_terms:
 return []

 terms = [item["term"] for item in candidate_terms]

 if method == "openai" and self.openai_available:
 try:
 # Get OpenAI embeddings
 chunk_embedding = self.openai_embeddings.embed_query(chunk)
 term_embeddings = self.openai_embeddings.embed_documents(terms)

 # Compute cosine similarities
 similarities = []
 chunk_emb = np.array(chunk_embedding)
 for term_emb in term_embeddings:
 term_emb = np.array(term_emb)
 similarity = np.dot(chunk_emb, term_emb) / (
 np.linalg.norm(chunk_emb) * np.linalg.norm(term_emb)
)
 similarities.append(float(similarity))
 except Exception as e:
 print(f"OpenAI embeddings error: {e}, falling back to sentence transformer")
 method = "sentence_transformer"

 if method == "sentence_transformer" or (method == "openai" and not self.openai_available):
 # Get sentence transformer embeddings
 all_texts = [chunk] + terms
 embeddings = self.sentence_model.encode(all_texts)
 chunk_emb = embeddings[0]
 term_embs = embeddings[1:]

 similarities = cosine_similarity([chunk_emb], term_embs)[0]
 similarities = [float(sim) for sim in similarities]

 elif method == "tfidf":
 # Use TF-IDF similarity
 all_texts = [chunk] + terms
 tfidf_matrix = self.tfidf_vectorizer.transform(all_texts)
 chunk_vector = tfidf_matrix[0]
 term_vectors = tfidf_matrix[1:]

 similarities = cosine_similarity(chunk_vector, term_vectors)[0]
 similarities = [float(sim) for sim in similarities]

 else:
 # Default to simple string similarity
 similarities = [0.5 for _ in terms]

 # Combine ensemble scores with similarity scores and generate rationale
 final_results = []
 for i, term_data in enumerate(candidate_terms):
 final_score = similarities[i]
 ensemble_score = term_data.get("ensemble_score", 0)

 # Weighted combination of ensemble retrieval and similarity
 combined_score = 0.6 * final_score + 0.4 * (ensemble_score / max(1, max(item.get("ensemble_score", 1) for i, item in enumerate(candidate_terms))))

 # Generate rationale for this match
 rationale = self._generate_match_rationale(
 chunk=chunk,
 term=term_data["term"],

```

```

 similarity_score=final_score,
 sources=term_data.get("sources", []),
 retrieval_method=method
)

 final_results.append({
 "term": term_data["term"],
 "similarity_score": final_score,
 "ensemble_score": ensemble_score,
 "combined_score": combined_score,
 "sources": term_data.get("sources", []),
 "retrieval_method": method,
 "rationale": rationale
 })

Sort by combined score
final_results.sort(key=lambda x: x["combined_score"], reverse=True)

return final_results

def search_chunk(self, chunk: str, similarity_method: str = "openai",
 top_k_retrieval: int = 15, top_k_final: int = 5,
 threshold: float = 0.7) -> Dict[str, Any]:
 """Complete search pipeline for a single chunk"""

 # Step 1: Ensemble retrieval to get candidate terms
 candidates = self.ensemble_retrieval(chunk, top_k=top_k_retrieval)

 if not candidates:

 return {
 "chunk_preview": chunk[:200] + "..." if len(chunk) > 200 else chunk,
 "candidates_found": 0,
 "final_matches": [],
 "status": "No candidates found"
 }

 # Step 2: Compute final similarities
 final_matches = self.compute_final_similarity(
 chunk, candidates, method=similarity_method
)[:top_k_final]

 # Step 3: Apply threshold
 threshold_matches = [
 match for match in final_matches
 if match["combined_score"] >= threshold
]

 best_match = threshold_matches[0] if threshold_matches else None

 return {
 "chunk_preview": chunk[:200] + "..." if len(chunk) > 200 else chunk,
 "candidates_found": len(candidates),
 "final_matches": final_matches,
 "threshold_matches": threshold_matches,
 "best_match": best_match,
 "status": "Language found" if best_match else "Not found"
 }

def search_multiple_chunks(self, chunks: List[str], similarity_method: str = "openai",
 top_k_retrieval: int = 15, top_k_final: int = 5,
 threshold: float = 0.7) -> List[Dict[str, Any]]:
 """Search multiple chunks efficiently"""

 results = []

 for i, chunk in enumerate(chunks):
 print(f"Processing chunk {i+1}/{len(chunks)}")

 result = self.search_chunk(
 chunk=chunk,

```

```

 similarity_method=similarity_method,
 top_k_retrieval=top_k_retrieval,
 top_k_final=top_k_final,
 threshold=threshold
)

 result["chunk_index"] = i
 results.append(result)

...
 return results

2. SMART SIMILARITY SEARCH (src/smart_similarity_search.py) - Updated
```python
from typing import List, Dict, Any
from src.ensemble_retriever import EnsembleRetriever
from src.pdf_processor import LangChainPDFProcessor
from src.langgraph_workflow import LanguageLookupWorkflow
import numpy as np
import json

class SmartSimilaritySearch:
    def __init__(self, languages_csv_path: str = "./data/languages.csv"):
        self.pdf_processor = LangChainPDFProcessor()
        self.ensemble_retriever = EnsembleRetriever(languages_csv_path)
        self.workflow = LanguageLookupWorkflow()

    def process_pdf_document(self, uploaded_file) -> Dict[str, Any]:
        """Process PDF document using LangChain"""
        try:
            # Process PDF into documents
            documents = self.pdf_processor.process_uploaded_pdf(uploaded_file)

            # Extract text chunks
            chunks = self.pdf_processor.process_text_chunks(documents)

            # Get enhanced chunk data
            enhanced_chunks = self.pdf_processor.create_enhanced_chunks(documents)

            return {
                "success": True,
                "documents": documents,
                "chunks": chunks,
                "enhanced_chunks": enhanced_chunks,
                "total_chunks": len(chunks),
                "metadata": {
                    "source_file": uploaded_file.name,
                    "total_pages": len(set(doc.metadata.get('page', 0) for doc in documents)),
                    "avg_chunk_size": np.mean([len(chunk) for chunk in chunks]) if chunks else 0
                }
            }
        except Exception as e:
            return {
                "success": False,
                "error": str(e),
                "documents": [],
                "chunks": [],
                "enhanced_chunks": [],

                "total_chunks": 0
            }

    def search_with_ensemble(self, chunks: List[str],
                            similarity_method: str = "openai",
                            top_k_retrieval: int = 15,
                            top_k_final: int = 5,
                            threshold: float = 0.7) -> List[Dict[str, Any]]:
        """Search using ensemble retrieval approach"""

        return self.ensemble_retriever.search_multiple_chunks(

```



```

        chunks=chunks,
        similarity_method=similarity_method,
        top_k_retrieval=top_k_retrieval,
        top_k_final=top_k_final,
        threshold=threshold
    )

def run_complete_analysis(self, uploaded_file,
                           similarity_method: str = "openai",
                           top_k_retrieval: int = 15,
                           top_k_final: int = 5,
                           threshold: float = 0.7,
                           enable_workflow: bool = True) -> Dict[str, Any]:
    """Run complete document analysis pipeline"""

    results = {
        "pdf_processing": {},
        "ensemble_search": {},
        "workflow_analysis": {},
        "summary": {}
    }

    # Step 1: Process PDF
    print("■ Processing PDF document...")
    pdf_results = self.process_pdf_document(uploaded_file)
    results["pdf_processing"] = pdf_results

    if not pdf_results["success"]:
        return results

    chunks = pdf_results["chunks"]

    # Step 2: Ensemble Search
    print("■ Running ensemble retrieval and similarity matching...")
    search_results = self.search_with_ensemble(
        chunks=chunks,
        similarity_method=similarity_method,
        top_k_retrieval=top_k_retrieval,
        top_k_final=top_k_final,

        threshold=threshold
    )

    # Calculate statistics
    found_matches = sum(1 for r in search_results if r["status"] == "Language found")
    success_rate = (found_matches / len(search_results) * 100) if search_results else 0

    results["ensemble_search"] = {
        "results": search_results,
        "total_chunks": len(chunks),
        "found_matches": found_matches,
        "success_rate": success_rate,
        "method_used": similarity_method
    }

    # Step 3: LangGraph Workflow (if enabled)
    if enable_workflow:
        print("■ Running LangGraph analysis workflow...")
        # Convert search results to format expected by workflow
        workflow_results = self.workflow.run_workflow(chunks, search_results)
        results["workflow_analysis"] = workflow_results

    # Step 4: Generate Summary
    results["summary"] = self._generate_summary(results)

    return results

def _generate_summary(self, results: Dict[str, Any]) -> Dict[str, Any]:
    """Generate summary statistics"""

    pdf_results = results.get("pdf_processing", {})

```

```

search_results = results.get("ensemble_search", {})
workflow_results = results.get("workflow_analysis", {})

summary = {
    "total_chunks": pdf_results.get("total_chunks", 0),
    "found_matches": search_results.get("found_matches", 0),
    "success_rate": search_results.get("success_rate", 0),
    "method_used": search_results.get("method_used", "unknown"),
    "workflow_enabled": bool(workflow_results),
    "processing_success": pdf_results.get("success", False),
    "search_success": bool(search_results.get("results")),
    "workflow_success": workflow_results.get("success", False) if workflow_results else None
}

# Add workflow insights if available
if workflow_results and workflow_results.get("success"):
    analysis = workflow_results.get("analysis", {})
    if "match_statistics" in analysis:
        summary.update(analysis["match_statistics"])

return summary

def export_results(self, results: Dict[str, Any], format: str = "json") -> str:
    """Export results in specified format with rationale included"""

    if format == "json":
        return json.dumps(results, indent=2, default=str)

    elif format == "markdown":
        summary = results.get("summary", {})
        search_results = results.get("ensemble_search", {}).get("results", [])
        workflow_report = results.get("workflow_analysis", {}).get("report", "")

        markdown_report = f"""
# Smart Ensemble Language Lookup Report

## Summary Statistics
- **Total Chunks**: {summary.get('total_chunks', 0)}
- **Found Matches**: {summary.get('found_matches', 0)}
- **Success Rate**: {summary.get('success_rate', 0):.1f}%
- **Method Used**: {summary.get('method_used', 'unknown')}

## Ensemble Retrieval Results
{#### Found Matches if summary.get('found_matches', 0) > 0 else #### No Matches Found}

"""

        for result in search_results[:10]: # Show first 10 results
            status_icon = "■" if result["status"] == "Language found" else "■"
            best_match = result.get("best_match", {})

            markdown_report += f"""
{status_icon} **Chunk {result.get('chunk_index', 0) + 1}**: {result['status']}
- **Text Preview**: {result['chunk_preview']}
- **Candidates Found**: {result.get('candidates_found', 0)}

"""

            if best_match:
                markdown_report += f"""- **Best Match**: {best_match.get('term', 'None')}
- **Combined Score**: {best_match.get('combined_score', 0):.3f}
- **Similarity Score**: {best_match.get('similarity_score', 0):.3f}
- **Ensemble Score**: {best_match.get('ensemble_score', 0):.3f}
- **Sources**: {' '.join(best_match.get('sources', []))}
- **Rationale**: {best_match.get('rationale', 'No rationale provided')}

"""

            markdown_report += "\n--\n"

        if workflow_report:

            markdown_report += f"\n## AI Analysis Report\n\n{workflow_report}"

```

```

        return markdown_report
    else:
        return str(results)
'''

## MATCH RATIONALE FEATURE EXPLANATION

### How the Rationale System Works:

1. **Direct Text Analysis**: Checks for exact mentions of language terms in the document chunk
2. **Partial Word Matching**: Identifies related terms and word components
3. **Semantic Similarity Assessment**: Categorizes the strength of semantic relationships
4. **Retrieval Method Context**: Explains which methods detected the match (LLM, BM25, TF-IDF)
5. **Embedding Analysis**: Describes the embedding method used for final similarity computation
6. **Technical Context Detection**: Identifies AI/ML related terms that provide supporting evidence

### Example Rationales:

- "Direct mention of 'Machine Learning' found in the text. Very high semantic similarity. Detection methods: LLM identification"
- "Related terms found: neural, network. High semantic similarity. Detection methods: Statistical text similarity detection"

### Benefits:

1. **Transparency**: Users understand why matches were made
2. **Confidence**: Higher confidence in system decisions
3. **Debugging**: Easier to identify false positives/negatives
4. **Education**: Learn about AI/ML relationships in documents
5. **Quality Assurance**: Validate matching logic

## INSTALLATION INSTRUCTIONS

1. Create project directory and virtual environment:
'''bash
mkdir langchain-lookup-project
cd langchain-lookup-project
python -m venv venv
source venv/bin/activate # On Windows: venv\Scripts\activate
'''

2. Install dependencies:
'''bash
pip install -r requirements.txt
'''

3. Set up environment variables:
'''bash

# Create .env file with your OpenAI API key
echo "OPENAI_API_KEY=your_openai_api_key_here" > .env
'''

4. Create data directory and add language database:
'''bash
mkdir data
# Add languages.csv file to data/ directory
'''

5. Run the Streamlit application:
'''bash
streamlit run app.py
'''

## USAGE

1. **Upload PDF**: Select a PDF document containing AI/ML content
2. **Configure Settings**: Choose similarity method, thresholds, and retrieval parameters
3. **Run Analysis**: Click "Analyze Document" to process the file
4. **Review Results**: Examine matches with detailed rationale explanations
5. **Export Data**: Download results as JSON or Markdown with reasoning included

```

The system now provides comprehensive explanations for every language match, making it transparent and trustworthy for

****System Architecture****: Ensemble Retrieval → Similarity Computation → Rationale Generation → Results Display

****Key Innovation****: Intelligent match explanation combining multiple evidence sources for transparent, explainable AI-p